Simple Masked Image Modeling

Major components

- Masking strategy - what to, how to mask
- Encoder - feature extraction, masked area prediction
- Prediction head - generative head
- Prediction target - target + loss

Masking Strategy

- replacing the masked patches w/ a mask token vector

- Patch-aligned random masking
  - the one described above, patch-level masking as it is convinient
- Other strategies have also been tried
  - central region masking
  - block-wise masking

Prediction head

↳ a single linear layer is sufficient

Prediction targets

↳ Raw pixel regression

↳ resized 1/x x32, x16, x8, x4, x2 down sampled targets

↳ L1 -loss (L2 and L1 smooth also tried)

Encoder

↳ Vanilla ViT and SWin transformer

(Same) Results

default patch:32 ratio:0,6

↳ Simple masking strategy seemed to be the best, with large patchsize (32) and masking % of 10 - 70

↓

on image size of 192^2

↳ larger masking ratio seem to help on smaller patchsizes as well

↳ smaller heads (1-2 linear layers) result in better fine-tuned results as the head can't learn much